

Stackable File System Using NFS Ganesha

Arun Olappamanna Vasudevan, *CS, Stony Brook University,*

Abstract—NFS Ganesha is a userland implementation of Network File System. This article explores stackable file system feature supported by NFS Ganesha version 2.0. A file-encrypting stackable file system CRYPTFS.

Keywords—*network file system, file system, file system abstraction layer, proxy, NFSv4, NFS, FSAL, stackable*

I. INTRODUCTION

GANESHA NFS being a userland implementation of Network File System protocol is a convenient place to try out new features and algorithms. Duo Yi and Youlong Cheng implemented security features in proxy server [1] using Ganesha NFS. The implementation is mostly in NFSv4 protocol, which makes it highly inflexible and non-modular. Stackable FSAL might help implement features in a cleaner way.

Stackable FSAL was explored by implementing a basic file-encrypting stackable file system in NFS Ganesha version 2.0.

II. PROJECT DESIGN

In order to design stackable file system in NFS Ganesha, an overview of NFS Ganesha is discussed first.

A. FSAL - File System Abstraction Layer

There are several types of file systems supported by Ganesha. The FSAL is an abstraction layer that abstracts each file system to a set of common operations that are handled by each module.

Each module can be configured using special key-value pairs in configuration file that is parsed by NFS Ganesha. For instance, FSAL_PROXY required Srv_Addr to be defined in configuration file.

1) **PROXY FSAL:** This is an abstraction that implements proxy machine between a client and a server. FSAL_PROXY uses NFSv4. Figure 1 shows how requests from clients are handled by FSAL_PROXY by contacting server.

- 1) Init module invokes config parser
- 2) 'proxy.ganesha.conf' is parsed

- 3) Library mentioned as value of 'FSAL_Shared_Library' is loaded
- 4) dlopen() invokes constructor that registers FSAL of library
- 5) FSAL added to fsal_list, call backs registered for configuration (init_config) and creating export (create_export)
- 6) EXPORT block in config file corresponds to a line in /etc/exports in Linux. Export entry is created with the FSAL mentioned in EXPORT block. In this case PROXY.
- 7) Proxy operations are registered with the export entry (function pointers for read, write, open, close, etc.)
- 8) After config file is parsed, all FSALs in fsal_list are initialized
- 9) Initialization of FSAL Proxy
- 10) Configuration items specific to FSAL Proxy, such as server address (Srv_addr) is obtained from config parse tree
- 11) Thread spawned to connect to server and initialize RPC
- 12) RPC socket that's used for send RPC request and listen
- 13) Worker threads spawned for handling requests from client
- 14) NFSv4 request from client received by worker thread
- 15) After access permission (supports even authentication) checks of client on export entry corresponding to request, NFSv4 service function (usually nfs4_Compound) is called
- 16) All metadata is maintained in cache inode and dentry data structures in AVL trees. Request received from NFSv4 operations.
- 17) Request is forwarded to FSAL operations registered to particular export entry
- 18) Server request through RPC socket
- 19) NFSv4 request to Server
- 20) NFSv4 response from Server
- 21) Response passed to PROXY handler
- 22) Response passed to cache inode
- 23) Response passed to NFSv4 protocol handler
- 24) Response passed to worker thread
- 25) svc_sendreply() to client

Arun O. V. is a graduate student with the Department of Computer Science, Stony Brook University, Stony Brook, NY, 11079 USA e-mail: aolappamanna@cs.stonybrook.edu, web: <http://www.fsl.cs.sunysb.edu/~arunov/>

